

Machine Learning


Gradient Descent

Tarapong Sreenuch


8 February 2024

克明峻德，格物致知

Recap: Gradients




$x \uparrow$ $y \uparrow$ gradient



$x \uparrow$ $y \downarrow$ gradient



Drawing a Line of Best Fit





Q: How do we find a line of best fit?

A: By rotating (clockwise/anti-clockwise) and/or shifting (up/down) the ruler, we find a line that goes roughly through the middle of all the middle of all the scatter plots.



Forming a Linear Model





Forming a Linear Model





Forming a Linear Model




Forming a Linear Model



Forming a Linear Model



Calculus



Q: Which value of x will $f(x)$ be either minimum or maximum?

Hint: What will happen to Slope (or Gradient) at those $x(s)$?

A: ... Slope (or Gradient) = 0 ...

Best Fit Line

Sum Squared Error (SSE):

$$SSE = (Y - X \times \vec{w})^T (Y - X \times \vec{w})$$


SSE Derivative:

$$\nabla SSE = 2X^T(Y - X \times \vec{w})$$


To minimise SSE, we find \vec{w} which results in $\nabla SSE = 0$.

Q: Which value of w will $f(w)$ be minimum?


A: ... Slope (or Gradient) = 0 ...



Non-Linear Equation: Gradients



Gradient Descent: Intuition




—ev gradient implies if w increased (move in the $+w$ direction) the SSE would also be decreased.

If SSE was to decrease, then we must move in the $+ w$ direction.

If SSE was continuing to decrease, then eventually we would reach the *minimum* point.

Hence, along as we have —ev gradient, for each step we are going to keep increasing w , i.e.
 $w^{(t+1)} = w^{(t)} + \Delta w$.

Gradient Descent: Intuition



+ev gradient implies if w increased (move in $+w$ direction) the SSE would be decreased.

If SSE was to decrease, then w must be decreasing, i.e. move in the $-w$ direction.

If w continued to decrease, then eventually SSE would reach the *minimum* point.

Hence, along as we have +ev gradient, for each step we are going to keep *decreasing* w , i.e.
 $w^{(t+1)} = w^{(t)} - \Delta w$.

Gradient Descent: Intuition

If +ve gradient, then $w^{(t+1)} = w^{(t)} - \Delta w$ else (-ev) $w^{(t+1)} = w^{(t)} + \Delta w$.




$$w^{(t+1)} = w^{(t)} - \text{sign}(\nabla \text{SSE}(w^{(t)})) \times \Delta w$$

gradient

Δw is fixed step. Unless is very small, w will *unlikely* be *very close* at the minimum point.

Gradient Descent: Observation




Wish List: We'd like Δw to be large when w is further away from the *minimum* point, and in the opposite smaller step when closer to the *minimum* point.

$|\nabla \text{SSE}(w)|$ becomes smaller when closer to the minimum point.

$$|\nabla \text{SSE}(w^{(t)})| \geq |\nabla \text{SSE}(w^{(t+1)})| \geq \dots \geq |\nabla \text{SSE}(w^{(t+N)})|$$

Gradient Descent: Revised



Pseudocode for Gradient Descent

```
# Inputs
# f(w)      ← objective to minimize (w → scalar)
# grad(w)   ← gradient of f at w (w → ∇f(w))
# w0        ← initial parameters
# n         ← learning rate (constant step size)
# max_iter  ← maximum number of iterations
# tol       ← stopping threshold on ||∇f(w)||

# ----- optimize -----
w ← w0

FOR t = 1 TO max_iter DO
    g ← grad(w)                  # compute gradient ∇f(w)

    IF norm(g) ≤ tol THEN        # stopping rule (first-order)
        BREAK
    END IF

    w ← w - n · g              # gradient descent update
END FOR

RETURN w
```

Gradient Descent from Scratch

```
from typing import Callable, Tuple
import numpy as np

def gradient_descent(
    grad: Callable[[np.ndarray], np.ndarray],
    w0: np.ndarray,
    eta: float = 0.05,
    max_iter: int = 1000,
    tol: float = 1e-6,
) -> Tuple[np.ndarray, int, float]:
    w = np.asarray(w0, dtype=float).ravel()
    n_iter = 0

    for t in range(max_iter):
        g = grad(w)
        gn = float(np.linalg.norm(g))
        if gn <= tol:
            n_iter = t
            break
        w = w - eta * g
        n_iter = t + 1
    else:
        # loop exhausted without break → recompute final grad norm
        gn = float(np.linalg.norm(grad(w)))

    return w, n_iter, gn
```

Linear Regression: Gradient Descent as a Minimiser

```
from sklearn.base import BaseEstimator, RegressorMixin
import numpy as np

class MyLinearRegressor(BaseEstimator, RegressorMixin):
    def __init__(self, eta=0.05, max_iter=1000, tol=1e-6):
        self.eta, self.max_iter, self.tol = eta, max_iter, tol

    def fit(self, X, y):
        # Design matrix with bias
        Phi = np.c_[np.ones(X.shape[0]), X]
        N, D = Phi.shape

        #  $\nabla \text{SSE}(w) = 2 \cdot \Phi^T(\Phi w - y)$ 
        def grad(w):
            r = Phi @ w - y
            return 2.0 * (Phi.T @ r)

        # Run the GD minimiser
        w0 = np.zeros(D)
        w, n_iter, gn = gradient_descent(
            grad, w0, eta=self.eta, max_iter=self.max_iter, tol=self.tol
        )

        self.weights_, self.n_iter_, self.grad_norm_ = w, n_iter, gn
        return self

    def predict(self, X):
        Xs = np.c_[np.ones(X.shape[0]), X]
        return Xs @ self.weights_
```


Usage Example

```
# gradient descent related settings
learning_rate = 0.01
max_iterations = 1000
tolerance      = 1e-6

# train with our sklearn-style class
model = MyLinearRegressor(eta=learning_rate, max_iter=max_iterations, tol=tolerance)
model.fit(X_train, y_train)

# optimised weights [bias, coefficients...]
print("Optimised Weights:", model.weights_)
```

```
Optimised Weights: [3.79008501 3.14520414]
```




Step 1: Weight Initialisation

```
# Step 1: Initialize weights  
w0 = np.zeros(D)
```

```
w, n_iter, gn = gradient_descent(  
    grad, w0, eta=self.eta, max_iter=self.max_iter,  
    tol=self.tol, callback=self.callback  
)
```


$$\vec{w} = \begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}$$

$$\begin{aligned}\hat{y} &= 0 \cdot x + 0 \\ &= 0\end{aligned}$$



Step 2: Gradient Descent Loop

```
# Step 2: Start the optimisation loop
for t in range(max_iter):
    g = grad(w)
    gn = float(np.linalg.norm(g))
    if gn <= tol:
        n_iter = t
        break
    w = w - eta * g
    n_iter = t + 1
    if callback is not None:
        callback(t, w, g, gn)
    else:
        # if not broken, recompute gn for reporting
        gn = float(np.linalg.norm(grad(w)))
```





Step 3: Return Optimised Weights


```
# Step 3: Return the optimized weights  
return w, n_iter, gn
```

```
# optimised weights [bias, coefficients...]  
print("Optimised Weights:", model.weights_)
```

```
Optimised Weights: [3.79008501 3.14520414]
```



Learning Rate




...just right...

- Low learning rate leads to slow convergence, which will require iterations.
- High learning rate can also lead to slow convergence, which is caused by oscillations.
- Very high learning rate will cause gradient descent not to converge.


If things go wrong, then just lower the learning rate.

High Learning Rate




Q: It looks like there are 2 trajectories, i.e. converges from bottom left and from top right. What happened?

A: ...Oscillation...



Workflow: Linear Regression



Summary

- Gradient Descent is an iterative optimization technique: It is widely used to minimize error functions by iteratively updating parameters in the direction of the negative gradient, aiming to reach the point where the function has the lowest error.
- Learning Rate determines the step size: The learning rate controls how much the parameters are adjusted with each iteration. Choosing an appropriate learning rate is crucial, as a high rate might cause overshooting, while a low rate could slow down convergence.
- Stopping Criteria ensure convergence: Gradient Descent uses stopping criteria to terminate the process once changes in the error fall below a specified threshold or a set number of iterations is reached, indicating that a minimum is close.
- Gradient Descent is in fact a the learning algorithm that iteratively adjusts model parameters to minimize error. Rather than being unique to any single model type (e.g., linear regression or neural networks), Gradient Descent is a general optimization technique used across various models to find the best parameter values that reduce the error metric.