


Machine Learning

Regularisation

Tarapong Sreenuch

8 February 2024

克明峻德，格物致知



I have a fever and running nose.

You're having a flu. I'll prescribe you paracetamol.


But, how about the virus (our cause)?

Q: How about overfitting in ML?

A: If we constrain its weights (the symptoms), then we can manage overfitting.




Really?

Example Dataset



Logistic Regression (2nd Order)

Feature	Weight
1	1.717
x_1	1.385
x_2	-0.579
x_1^2	-0.167
x_2^2	0.978




Logistic Regression (8th Order)


Feature	Weight
1	587.984
x_1	495.169
x_2	-492.094
x_1^2	-76.356
x_2^2	-409.916
x_1^3	-53.310
x_2^3	286.153
x_1^4	-356.961
x_2^4	-120.300
x_1^5	-457.228
x_2^5	540.862

Feature	Weight
x_1^6	74.364
x_2^6	233.336
x_1^7	166.729
x_2^7	-167.182
x_1^8	17.001
x_2^8	-31.508

Overfitting → Large Weights



Workflow: Logistic Regression



Desired Total Cost

Want to Balance

$$\text{Total Cost} = \text{Measure of Fit} + \text{Measure of Magnitude of Weights}$$

$$-L(\vec{w})$$

Negative Log Likelihood

$$\|\vec{w}\|_2^2 = w_1^2 + \dots + w_M^2$$

Resulting Objective

Minimising Total Cost Function:

$$\min_{\vec{w}} -L(\vec{w}) + \lambda \|\vec{w}\|_2^2$$

Penalty Parameter, i.e. Balance of Fit and Magnitude

Ridge Regression (aka. L2 Regularisation)

Gradient Vector

Total Cost Function:


$$J(\vec{w}) = - \sum_{i=1}^N [y^{(i)} \log P(1|\vec{x}^{(i)}, \vec{w}) + (1 - y^{(i)}) \log(1 - P(1|\vec{x}^{(i)}, \vec{w}))] + \lambda \|\vec{w}\|_2^2$$

Derivative of Total Cost Function:

$$\frac{\partial J(\vec{w})}{\partial w_j} = - \sum_{i=1}^N (y_i - P(1|\vec{x}^{(i)}, \vec{w})) x_j^{(i)} + 2\lambda \vec{w}$$

Gradient Vector:

$$\nabla_{\vec{w}} J(\vec{w}) = -X^T(Y - P(1|X, \vec{w})) + 2\lambda \vec{w}$$



Pseudocode for L_2 Regularisation

```
# L2 Regularisation – Plug-in Gradient Term

# Inputs
# base_grad(w) ← gradient of the base loss (no regularisation)
# w0           ← initial parameters
# n            ← learning rate
# max_iter     ← maximum iterations
# tol          ← stop when  $\|g\| \leq \text{tol}$ 
#  $\lambda \geq 0$     ← L2 strength (intercept  $w[0]$  is never penalised)

# ----- helper (exclude intercept) -----
Function REG_VECTOR(w)
    r ← copy(w)
    r[0] ← 0
    Return r
End

# L2 loss piece:  $\lambda \cdot \|\text{REG\_VECTOR}(w)\|^2$ 
# L2 grad piece:  $2 \cdot \lambda \cdot \text{REG\_VECTOR}(w)$ 

# ----- optimise (vanilla GD + L2) -----
w ← w0

FOR t = 1 TO max_iter DO
    g_base ← base_grad(w)          # gradient of base loss
    g_reg  ← 2 · λ · REG_VECTOR(w) # L2 gradient (no intercept)
    g      ← g_base + g_reg       # total gradient

    IF norm(g) ≤ tol THEN BREAK
    w ← w - n · g
END FOR

Return w
```

Logistic Regression with L_2 Regularisation

```
from sklearn.base import BaseEstimator, ClassifierMixin
import numpy as np
from scipy.special import expit # stable sigmoid

class MyLogisticRegression(BaseEstimator, ClassifierMixin):
    def __init__(self, eta=1e-3, max_iter=2000, tol=1e-6, lambda_=0.0):
        self.eta, self.max_iter, self.tol = eta, max_iter, tol
        self.lambda_ = lambda_ # λ: L2 strength (0.0 = no regularization)

    def fit(self, X, y):
        Phi = np.c_[np.ones(X.shape[0]), X] # [1 | X]
        N, D = Phi.shape

        # Gradient of: sum NLL + λ*||w'||^2 (w'[0]=0 ⇒ intercept not penalized)
        def grad(w):
            p = expit(Phi @ w)
            reg = w.copy(); reg[0] = 0.0
            return (Phi.T @ (p - y)) + self.lambda_ * reg

        w0 = np.zeros(D)
        w, n_iter, gn = gradient_descent(
            grad, w0, eta=self.eta, max_iter=self.max_iter, tol=self.tol
        )

        self.weights_, self.n_iter_, self.grad_norm_ = w, n_iter, gn
        return self

    def predict_proba(self, X):
        Xs = np.c_[np.ones(X.shape[0]), X]
        p1 = expit(Xs @ self.weights_)
        return np.column_stack([1 - p1, p1])

    def predict(self, X):
        return (self.predict_proba(X)[:, 1] ≥ 0.5).astype(int)
```

Gradient Descent Function

```
from typing import Callable, Tuple
import numpy as np

def gradient_descent(
    grad: Callable[[np.ndarray], np.ndarray],
    w0: np.ndarray,
    eta: float = 0.05,
    max_iter: int = 1000,
    tol: float = 1e-6,
) → Tuple[np.ndarray, int, float]:
    w = np.asarray(w0, dtype=float).ravel()
    n_iter = 0

    for t in range(max_iter):
        g = grad(w)
        gn = float(np.linalg.norm(g))
        if gn ≤ tol:
            n_iter = t
            break
        w = w - eta * g
        n_iter = t + 1
    else:
        # loop exhausted without break → recompute final grad norm
        gn = float(np.linalg.norm(grad(w)))

    return w, n_iter, gn
```

Usage Example

```
import numpy as np
# from your_module import MyLogisticRegression # class already defined

# ----- data -----
m, n = 100, 2
np.random.seed(0)
class_0 = np.hstack((1.5 + np.random.randn(m, 1), -1.5 + np.random.randn(m, 1)))
class_1 = np.hstack((-1.5 + np.random.randn(m, 1), 1.5 + np.random.randn(m, 1)))
X = np.vstack((class_0, class_1)) # (2m, 2)
y = np.concatenate([np.zeros(m), np.ones(m)]) # (2m,) in {0,1}


# test point (no bias term; model adds it)
X_test = np.array([[1.0, -1.0]])

# ----- train with L2 only (use default eta/max_iter/tol) -----
lambda_ = 0.1
model = MyLogisticRegression(lambda_=lambda_) # default GD param
model.fit(X, y)


# ----- predict -----
prob = model.predict_proba(X_test)[0, 1]

print(f"Predicted probability for new point: {prob:.4f}")
```

Predicted probability for new point: [0.1415866]




Example Dataset



Train, Validation and Test Datasets

```
from sklearn.model_selection import train_test_split  
  
# First, split the data into train (70%) and test (30%)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
```




Train

Test

- Train:Test is typically either 0.8:0.2 or 0.7:0.3.
- Test dataset is a proxy of unseen data, and it will only be used in the final evaluation.
- Train dataset is further divided into k folds (e.g., 5 or 10). For each fold, the model is trained on $k-1$ folds and validated on the remaining fold.
- We use **K-Fold Cross-Validation** to fine-tune or optimize the ML model. Here, we find a hyper-parameter λ based on the average performance across folds.

K-Fold Cross Validation




For $k = 1, \dots, K$

Step 1: Estimate $\vec{w}_\lambda^{(k)}$ on the training blocks.

Step 2: Compute error on Validation Block: $\text{Error}^{(k)}(\lambda)$.

K-Fold Cross Validation




For $k = 1, \dots, K$

Step 1: Estimate $\vec{w}_\lambda^{(k)}$ on the training blocks.

Step 2: Compute error on Validation Block: $\text{Error}^{(k)}(\lambda)$.

K-Fold Cross Validation




For $k = 1, \dots, K$

Step 1: Estimate $\vec{w}_\lambda^{(k)}$ on the training blocks.

Step 2: Compute error on Validation Block: $\text{Error}^{(k)}(\lambda)$.

K-Fold Cross Validation




For $k = 1, \dots, K$

Step 1: Estimate $\vec{w}_\lambda^{(k)}$ on the training blocks.

Step 2: Compute error on Validation Block: $\text{Error}^{(k)}(\lambda)$.

K-Fold Cross Validation



For $k = 1, \dots, K$

Step 1: Estimate $\vec{w}_\lambda^{(k)}$ on the training blocks.

Step 2: Compute error on Validation Block: $\text{Error}^{(1)}(\lambda)$.

Compute Average Error: $\text{CV}(\lambda) = \frac{1}{K} \sum_{k=1}^K \text{Error}^{(k)}(\lambda)$.

scikit-learn: Logistic Regression with L_2 Regularisation

```
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold, cross_val_score

# Step 1: Create an instance of LogisticRegression
lambda_ = 10000
model = LogisticRegression(penalty='l2', C=1/lambda_, random_state=204)

# Step 2: Define cross-validation strategy
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Step 3: Perform cross-validation on the training data
cv_scores = cross_val_score(model, X_train, y_train, cv=kf, scoring='accuracy')
print("Mean Cross-Validation Accuracy: {:.2f}".format(np.mean(cv_scores)))
# optional: print per-fold scores
# print("Per-fold:", cv_scores)
```




Mean Cross-Validation Accuracy: 0.77







Regularised Decision Boundary (8th Order)

Feature	Coefficient
1	0.609
x_1	3.661e-4
x_2	-3.319e-5
x_1^2	-6.429e-5
x_2^2	-2.815e-5
x_1^3	8.919e-4
x_2^3	-1.843e-4
x_1^4	-1.316e-4
x_2^4	-2.531e-4
x_1^5	2.826e-3
x_2^5	-8.357e-4


Feature	Coefficient
x_1^6	-2.408e-4
x_2^6	-1.386e-3
x_1^7	9.870e-3
x_2^7	-3.375e-3
x_1^8	-3.631e-4
x_2^8	-6.624e-3





Model Complexity



Bias-Variance Trade-Offs



- High Variance, Low Bias
- Complex Decision Boundary
- High Train Accuracy
- Low Validation Accuracy




- High Bias, Low Variance
- Simple Decision Boundary
- Low Train Accuracy
- Low Validation Accuracy

Regularised Decision Boundary ($\lambda = 10^{13}$)

Feature	Coefficient
1	-0.194
x_1	1.463e-12
x_2	-5.507e-13
x_1^2	-2.369e-12
x_2^2	-1.633e-12
x_1^3	1.310e-11
x_2^3	-4.208e-12
x_1^4	-4.208e-11
x_2^4	-1.171e-11
x_1^5	1.935e-10
x_2^5	-3.018e-11


Feature	Coefficient
x_1^6	-7.800e-10
x_2^6	-8.390e-11
x_1^7	3.483e-9
x_2^7	-2.244e-10
x_1^8	1.508e-8
x_2^8	-6.308e-10



Regularised Decision Boundary ($\lambda = 10^{-7}$)


Feature	Coefficient
1	432.548
x_1	333.613
x_2	-352.095
x_1^2	-61.796
x_2^2	-256.131
x_1^3	-36.690
x_2^3	219.2734
x_1^4	-246.131
x_2^4	-106.567
x_1^5	-310.860
x_2^5	405.405

Feature	Coefficient
x_1^6	44.591
x_2^6	107.260
x_1^7	110.216
x_2^7	-150.206
x_1^8	11.905
x_2^8	15.687




Overfitting vs Underfitting


- **Overfitting:** With minimal penalty on weights, the model captures noise in the training data, leading to high variance and poor generalization on validation data.
- **Underfitting:** A strong penalty on weights forces them to shrink, overly simplifying the model. This results in high bias, causing it to miss important patterns in both training and validation data.



- High Variance, Low Bias
- Complex Decision Boundary
- High Train Accuracy
- Low Validation Accuracy



- **Optimal:** The ideal λ imposes just enough weight penalty to prevent overfitting while preserving the model's ability to capture true patterns, leading to accurate predictions on both training and unseen data.



- High Bias, Low Variance
- Simple Decision Boundary
- Low Train Accuracy
- Low Validation Accuracy

Summary

- Regularisation helps prevent overfitting by penalising large weights, encouraging models to capture meaningful patterns rather than noise. This leads to simpler, more generalised models that perform well on new data.
- The cost function in regularised logistic regression combines the negative log-likelihood with an added penalty term, $\lambda \|\vec{w}\|_2^2$. This term balances minimising error with controlling weight magnitudes to reduce model complexity.
- The regularisation parameter, λ , governs the strength of the penalty. A smaller λ may result in high variance and overfitting, while a larger λ increases bias and risks underfitting. Tuning λ helps strike an optimal balance between variance and bias.
- Regularisation implicitly controls model complexity and performs feature selection by suppressing coefficients. Features with low-magnitude weights are essentially excluded from the model, allowing it to focus on the most impactful features and reducing overfitting.